

MACHINE LEARNING

Spring 2017

Contact: Patrick Beukema
Email: plb23@pitt.edu

Time: W 12:00 – 1:00
Place: Mellon Institute 115

Course Page:

1. Course website: http://github.com/pbeukema/ML_CNBC
2. Code referenced in text: <http://github.com/probml/pmtk1/tree/master/pmtk/bookCode>

Main Text: Machine Learning: A Probabilistic Perspective. <https://www.cs.ubc.ca/~murphyk/MLbook/>. Each week we will read, do exercises, and discuss one chapter of this text.

Schedule & Exercises:

1. Introduction 1/11
 - Excercises: 1.1-1.3, Rec: Write KNN from scratch
2. Probability 1/18
 - Exercises: 2.1-2.5, 2.12, Rec: 2.17 (Prove & Run simulation)
3. Generative models for discrete data 1/25
 - Exercises
4. Gaussian models 2/1
 - Exercises
5. Bayesian statistics 2/8
 - Exercises
6. Frequentist statistics 2/15
 - Exercises
7. Linear Regression 2/22
 - Exercises
8. Logistic Regression 3/1
 - Exercises
9. Generalized linear models and the exponential family 3/8
 - Exercises
10. Directed graphical models 3/15
 - Exercises

11. Mixture models and the EM algorithm 3/22

- Exercises

12. Latent linear models 3/29

- Exercises

13. Sparse linear models 4/5

- Exercises

14. Kernels 4/12

- Exercises

15. Gaussian processes 4/19

- Exercises

Course objectives: The idea behind this course is to gain exposure and learn the fundamentals of machine learning through weekly discussions and exercises. Secondary goal is to republish the code that is used throughout the text and excersises in the form of jupyter notebooks in python. The motivation is for that is twofold. 1. Current implemntations are in Matlab, which some people lack access to. 2. More learning will happen doing things from scratch.

Prerequisites: It will be challenging to follow along without some bacground in calculus, linear algebra, probability and statistics.

Resources & Recommendations:

- Learning python: learnpythonthehardway.org
- Programming: jupyter.org notebooks
- Editor: atom.io