

MACHINE LEARNING

Spring 2017

Contact: Patrick Beukema
Email: plb23@pitt.edu

Time: W 12:00 – 1:00
Place: Mellon Institute 115

Course Page: http://github.com/pbeukema/ML_CNBC

Main Text: Each Week we will read and do exercises from one chapter of Machine Learning: A Probabilistic Perspective. <https://www.cs.ubc.ca/~murphyk/MLbook/>

Schedule & Exercises:

1. Introduction
 - Excercises: 1.1-1.3, Rec: Write KNN from scratch
2. Probability
 - Excercises: 2.1-2.5, 2.12, Rec: 2.17 (Prove & Run simulation)
3. Generative models for discrete data
 - Excercises
4. Gaussian models
 - Excercises
5. Bayesian statistics
 - Excercises
6. Frequentist statistics
 - Excercises
7. Linear Regression
 - Excercises
8. Logistic Regression
 - Excercises
9. Generalized linear models and the exponential family
 - Excercises
10. Directed graphical models
 - Excercises
11. Mixture models and the EM algorithm
 - Excercises

12. Latent linear models

- Excercises

13. Sparse linear models

- Excercises

14. Kernels

- Excercises

15. Gaussian processes

- Excercises

Course objectives: This course is primarily designed for graduate students, and will introduce an audience to the state-of-the-art in modeling techniques for computer science and engineering majors. We try to discuss as many models as possible. We chiefly focus on complex networks, inference, machine learning, and probabilistic/statistical models and methods.

Prerequisites: It will be challenging to follow along without some background in calculus, linear algebra, probability and statistics.

Important Dates:

Midterm #1 Ābān 16, 1393 \equiv November 7, 2014
Midterm #2 Āzar 21, 1393 \equiv December 12, 2014
Final Exam Dey 18, 1393 \equiv January 8, 2015

Resources & Recommendations:

- Learning python: Learn Python the Hard Way learnpythonthehardway.org/
- Programming: Jupyter notebooks jupyter.org
- Text editor (& Programming): Atom atom.io
- Github: